

## 0 Contents

<b>1</b>	<b>C</b>	<b>2</b>
1.1	fixed size types . . . . .	2
1.2	Addition of pointers . . . . .	2
1.3	Index Operator on Pointers . . . . .	3
1.4	Padding . . . . .	3
1.5	Forwards Declaration . . . . .	3
<b>2</b>	<b>Operating System APIs</b>	<b>3</b>
2.1	Basic features of an Operating system . . . . .	3
2.1.1	Limitations of Portability . . . . .	3
2.1.2	Limits of Isolation . . . . .	3
2.2	Processor Privilege . . . . .	3
2.3	Kernel . . . . .	3
2.3.1	MicroKernel . . . . .	3
2.3.2	Monolithic Kernel . . . . .	3
2.3.3	Unikernel . . . . .	4
2.3.4	Running an instruction in Kernel Mode . . . . .	4
2.3.5	Syscall (SVC on ARM) . . . . .	4
2.4	ABI vs API . . . . .	4
2.4.1	ABI in Linux . . . . .	4
2.4.2	API in Linux . . . . .	4
2.5	POSIX . . . . .	4
2.5.1	POSIX Conformity . . . . .	4
2.6	Man Pages . . . . .	4
2.7	Shell . . . . .	4
2.7.1	Arguments in shell . . . . .	4
2.7.2	Env Vars . . . . .	4
<b>3</b>	<b>Filesystems</b>	<b>5</b>
3.1	Files . . . . .	5
3.1.1	FileTypes . . . . .	5
3.2	Directories . . . . .	5
3.2.1	Special directories . . . . .	5
3.2.2	Paths . . . . .	5
3.2.3	Max Path . . . . .	5
3.2.4	Rights . . . . .	6
<b>4</b>	<b>File API</b>	<b>6</b>
4.1	POSIX File API . . . . .	6
4.1.1	Usage of errno . . . . .	6
4.1.2	File-Descriptor . . . . .	6
4.1.3	Opening files with POSIX API . . . . .	6
4.1.4	Close files with POSIX API . . . . .	6
4.1.5	Usage of open and close . . . . .	7
4.1.6	Read data with POSIX API . . . . .	7
4.1.7	Write data with POSIX API . . . . .	7
4.1.8	Jump in a file with POSIX API . . . . .	7
4.1.9	pread and pwrite in POSIX API . . . . .	7
4.1.10	windooof proprietary paths . . . . .	7
4.1.11	Example for reading and writing data in POSIX API . . . . .	8
4.2	C Stream API . . . . .	8
4.2.1	FILE datastructure . . . . .	8
4.2.2	open file with C-API . . . . .	8
4.2.3	close file with C-API . . . . .	8
4.2.4	flush of file with C-API . . . . .	8
4.2.5	Conversion from POSIX-API to C-API . . . . .	8
4.2.6	read from file with C-API . . . . .	8
4.2.7	"un"read from file with C-API . . . . .	9
4.2.8	Write to file with C-API . . . . .	9
4.2.9	End of file and Error in File C-API . . . . .	9
4.2.10	Manipulation of file-position indicator with C-API . . . . .	9
4.3	Ext2 . . . . .	9
4.4	Ext4 . . . . .	9
<b>5</b>	<b>Processmodels</b>	<b>9</b>
5.1	Base . . . . .	9
5.1.1	Improvement with "async" . . . . .	10
5.2	Process . . . . .	10
5.2.1	Program vs Process . . . . .	10
5.2.2	Process Control Block (PCB) . . . . .	10
5.2.3	Interrupts, Processes and Context Switch . . . . .	10
5.2.4	Creation of a process . . . . .	10
5.3	Process Hierarchy . . . . .	10
5.3.1	fork function . . . . .	10
5.3.2	exit function . . . . .	11
5.3.3	wait function . . . . .	11
5.3.4	waitpid function . . . . .	11
5.3.5	Examples . . . . .	11
5.3.6	exec functions . . . . .	11
5.3.7	Zombie Processe . . . . .	12
5.3.8	Orphan Process . . . . .	12
5.3.9	function sleep . . . . .	12
5.3.10	function atexit . . . . .	12
5.3.11	function to read pid . . . . .	12

<b>6</b>	<b>C Toolchain</b>	<b>12</b>
6.1	Precprocessor	12
6.2	Linker	12
6.3	Loader	13
6.3.1	Penguin Loader	13
6.4	Executable Linking Format ELF	13
6.4.1	ELF Structure	13
6.4.2	Segments and Sections	13
6.4.3	Header of an ELF file	14
6.4.4	Program Header Table and Segments	14
6.4.5	Sections	14
6.4.6	String Table	14
6.4.7	Symbols	15
6.5	Static Libraries	15
6.6	Dynamic Libraries	15
6.6.1	Delayed Loading	15
6.7	POSIX API: dynamic libraries	16
6.7.1	dlopen	16
6.7.2	dlsym	16
6.7.3	dlclose	16
6.7.4	dlerror	16
6.7.5	Automatic Loading of ELF files	16
6.8	Naming of Shared Objects	16
6.8.1	Updates	16
6.9	Shared Objects with Linker and Loader	16
6.10	Creating Static Libraries:	17
6.11	Creating Dynamic Libraries:	17
6.12	Using libraries	17
6.12.1	ld-linux.so effective loader	17
6.13	Shared Object Facts	17
6.14	Dynamic Library Implementation	17
6.14.1	Position-Dependent Code	18
6.14.2	Position-Independent Code	18
6.14.3	Global Offset Table	18
6.15	PLT Process Linkage Table	18
<b>7</b>	<b>MultiThreading</b>	<b>18</b>
7.1	Threads	18
7.2	Process-Model vs Thread-Model	19
7.3	Thread as stack and context	19
7.4	Parallization of algorithms	19
7.4.1	Time difference	19
7.4.2	Speedup Factor	19
7.5	Amdahls Rule	20
<b>8</b>	<b>POSIX Threading API</b>	<b>20</b>
8.1	Creating a thread	20
8.1.1	Example for pthread	20
8.1.2	Attributes for pthread	20
8.2	pthread_exit	20
8.2.1	Lifetime of a thread	21
8.3	pthread_cancel	21
8.4	pthread_detach	21
8.5	pthread_join	21
8.6	pthread_self	21
<b>9</b>	<b>Thread-Local Storage (TLS)</b>	<b>21</b>
9.1	pthread_key_create	21
9.2	pthread_key_delete	21
9.3	pthread_setspecific and pthread_getspecific	21
9.4	Example with TLS	21
<b>10</b>	<b>Communication and Synchronization</b>	<b>22</b>
<b>11</b>	<b>Programs and libraries</b>	<b>22</b>
<b>12</b>	<b>Graphical Overlays</b>	<b>22</b>

## 1 C

### 1.1 fixed size types

- `int8_t`, `int16_t`, `int32_t`, `int64_t`  
fixed integers with bit count
- `intmax_t` max size int on platform
- `intptr_t` signed integer with the size of an address on this platform
- `uint8_t`, `uintptr_t` unsigned versions
- `size_t`  
this is used in containers, the reason for this is that *this has the max size that for example an array can be*.  
**This is unsigned!**

### 1.2 Addition of pointers

If you try to add or subtract 2 pointers to get the amount of `sizeof(t)` difference, then you can only do this with the exact same type, something like signed int and unsigned int will not work!

```
int32_t *y = 100;
int32_t *x = 120;
ptrdiff_t z = x - y; // z == 5
uint32_t *u = 120;
ptrdiff_t p = u - y; // Error: Different ptr types
```

### 1.3 Index Operator on Pointers

You can index on pointers like an array, this can be used to get elements on any object.

Note that you have to manually make sure to stay within the bounds of that object, as otherwise you will have *undefined behavior*.

```
int32_t x = 0;
int32_t *y = &x;
y[0] = 0x42;      // same as: x = 0x42;
(&x)[0] = 0x42;    // same
0[&x] = 0x42;      // same
100[200] = 0x42;   // Error: no address
```

### 1.4 Padding

When you mix and match different types of different sizes inside of a struct, then the compiler will include padding based on the bigger type:

```
struct {
char c;      // Offset 0
int32_t x;   // Offset 4 --> Padding
char d;      // Offset 8
} t;         // sizeof t == 12
# structure matters!!
struct {
char c;      // Offset 0
char d;      // Offset 1
int32_t x;   // Offset 2 --> Padding
} t;         // sizeof t == 6
```

### 1.5 Forwards Declaration

```
struct Folder;
// Forward-Deklaration
struct File {
struct Folder *parent;
// OK: all pointer types
//
have same size
char name[256];
// OK: fixed size array
};
// --> Type complete
struct Folder {
struct File * file[256]; // OK: fixed size array
};
// --> Type complete
```

## 2 Operating System APIs

### 2.1 Basic features of an Operating system

- **abstraction and portability**
  - define generic APIs that work on all (as many as possible) devices
  - define abstractions that we don't care about -> how are files stored on the disk?
- **Isolation and Resource Management**
  - Isolate each usecase from each other -> posix
  - Runtime (make it blazingly fast)
  - Memory Management
  - Secondary Storage handling
- **Security**

#### 2.1.1 Limitations of Portability

While the operating system can define standards, there are often things that we as developers need to consider, for example, while the operating system can define how a user will interact with the keyboard etc, if said device *doesn't have this input*, then your application will not work... Eg. An application meant for touch on the desktop might work, but not properly, and the OS can't really help there.

#### 2.1.2 Limits of Isolation

Often, you want some form of interoperability, or you are basically forced to use that.

For example an application might want the focus of the keyboard, but then a popup appears.

If the application continues taking the focus, then the application now breaks the user experience.

### 2.2 Processor Privilege

Modern operating systems define a range of instructions that only the kernel is allowed to perform.

This is done to protect the operating system from attacks that might be exploitable via these privileges.

In this case you run in *user mode*, this is also why anti-cheats are often running in kernel mode,

in this mode, the anti-cheat can access any memory all the time for whatever reason the anti-cheat would like to do so.

Should an application break the rule of user-space, the operating system will be notified and can then kill or otherwise restrict the application.

### 2.3 Kernel

#### 2.3.1 MicroKernel

Microkernel is the idea that only the *absolutely necessary operations need to be in the kernel*, this means that often, things like drivers run in the userspace, eg. Radv would be in userspace.

Features:

- **Reliability**  
less code is easier to maintain, which means a more stable operating system.
  - **Analysability**  
less code is easier to bisect, meaning that bug hunting is easier.
  - **Performance**  
Because drivers are now in a lower priority environment, they can no longer directly access hardware.  
This means that you will have a significant performance hit, which is also the reason that *linux is not a microkernel!*
- In the real world, there is no *real microkernel*, they usually add the necessary functionality of drivers and leave it at that.

#### 2.3.2 Monolithic Kernel

Monolithic kernels have all the base functionality included. This means that you will not need to supply basic functionality to the kernel, just to get a functional operating system.

- **Performance**  
Since drivers have direct access to hardware, this means they can run faster!

- **Security**  
Since drivers have direct access to hardware, this means that misconfigured or malicious hardware, can easily infiltrate the kernel
- **Reliability**  
More code means more possible bugs, and in the kernel this is worse than in the usermode.

### 2.3.3 Unikernel

This is a kernel that is made for one specific application, which means *it is an application!*.

- Performance
- Security
- Reliability
- Only one Usecase

### 2.3.4 Running an instruction in Kernel Mode

When you want to run an instruction in the kernel mode, then you need to do a *syscall*.

The processor will then switch into kernel mode (if the os has given the privilege) and run the instruction in kernel mode.

### 2.3.5 Syscall (SVC on ARM)

There is only one function to run something in kernel mode, this means that we have to use *codes* instead.

Eg. a syscall with code 60 would be the exit code for a program. -> plox kill me

**NOTE:** Syscall also doesn't take arguments, therefore you need to place the arguments in registers.

This is exactly why you had to place all these things into registers, when you wanted to print a simple "hello world" in assembly.

*The implication: output and input are kernel mode!!*

## 2.4 ABI vs API

Application Binary Interface

- concrete interfaces
- calling conventions
- projection of datastructures

Application Programming Interface

- abstract interfaces
- platform/OS independent aspects

### 2.4.1 ABI in Linux

Calling Conventions for syscall are different for different linux kernels!

This means that you need to compile applications for each kernel!

To counter problems that will appear with this, there is a standard called *Linux Standard Base*, which defines a set of conventions to use.

### 2.4.2 API in Linux

The proper solution is to use APIs instead, which can be done with languages such as C (and tomorrow Rust).

This means that you *no longer use syscall*, you instead use *C functions*, which work on every kernel, not just on one.

## 2.5 POSIX

In general, every OS has its own ABI and API.

The unix API has been developed alongside the C API, this lead to the ISO standard.

However, at some point there were multiple standards, which meant the compatability was wrecked again.

Instead, the POSIX standard API was defined, which meant that if you wrote your program POSIX compliant, then it will run on any POSIX OS.

### 2.5.1 POSIX Conformity

- MacOS: since version 10.5
- Linux, not certified, but somewhat POSIX conform  
Bad: not everything is standard, but we all know that sometimes you either go your own way, or nothing happens -> matrix
- Windows: no :)
- BSD: yes

## 2.6 Man Pages

Man pages provide information about a POSIX system, it *is made of 9 parts*:

1. Executable Programs or shell commands
2. System calls (functions provided by the kernel)
3. Library calls (functions within program libraries)
4. Special files (usually found in /dev)
5. File formats and conventions, e.g. /etc/passwd
6. Games lol
7. Miscellaneous (including macro packages and conventions)
8. System administration commands (usually only for root)
9. Kernel routines (not standard)

## 2.7 Shell

- Made of an output and input stream.
- many different shells, bash, dash, zsh, fish, nu
- doesn't need special rights or prerequisites
- Made to call OS functions via text

### 2.7.1 Arguments in shell

- All arguments are considered strings -> its just an IO stream
- Spaces usually separate the arguments
- "\" usually used to escape characters -> like space

These arguments are then passed to C or C++ program in the array called (*char \*\*argv*), and the (*int argc*) variable is the count of arguments.

**The first entry in argv is the programname**

### 2.7.2 Env Vars

- Key-Value pair
- Example: MOZ\_ENABLE\_WAYLAND=1
- can be set for the shell in .zshrc/.bashrc etc
- can be set for environment in xdg-config/environment.d, /etc/environment or .profile, /etc/profile
- Order is root first then local, if variable is set twice, then you will overwrite it!

**To use variables in C: getenv, putenv, setenv, unsetenv**

Technically, there is a variable called *environ*, which is an array of null pointers to strings which is 0 terminated.

However, this variable *is not defined!!*

**getenv**

```
//definition
char * getenv (const char * key)

char *value = getenv("PATH");
```

```
// value = "/home/ost/bin:/home/ost/.local/bin"
// returns nullptr -> 0 if variable not set
```

#### setenv

```
// definition
int setenv(const char *key, const char *value, int overwrite);

int ret = setenv("HOME", "/usr/home", 1);
// returns code 0 if ok, error code otherwise
// sets variable with value, overwrite if not 0
// error if variable doesn't exist!
```

#### unsetenv

```
// definition
int unsetenv(const char *key);

int ret = unsetenv("HOME");
// returns 0 if ok, error code otherwise
// removes the env variable
```

#### putenv

```
//definition
int putenv (char * kvp)

int ret = putenv("HOME=/usr/home");
// adds env variable pair
// returns 0 if ok, error code otherwise
// if variable already exists -> overwrite
```

In general, use env vars as a flag to configure things, not as a config that you *need* to configure.

For other operating systems this is done via other management solutions like windows -> registry.... haah get shit on windoof users

### 3 Filesystems

Essentially an API that makes sure that applications do not need to understand/know how the hardware works

#### 3.1 Files

Files have 2 parts:

- **Data**  
Sequence of bytes that represent the file
- **Metadata**  
visible for users: size, date, owner, filetype hint  
not visible for users: place on hardware, connection of blocks on hardware

##### 3.1.1 FileTypes

- ending after . -> .pdf
- File-endings have pretty much no point, they are just for the user
- File-endings are used as a hint to open said file with a specific program
- Type can be deducted via *magic numbers* inside of the file

General Advice:

- Data is trash unless proven otherwise
- Validate ALL data
- General advice: when decoding a file, continuously check if the file is really a pdf, or whatever you expect, test it to be.  
A few lines may not be enough to prove that it really is of said type.

#### 3.2 Directories

- Essentially a file with a special type
- Each directory other than root directory has exactly one parent -> tree
- root directory is often / -> penguinOS

##### 3.2.1 Special directories

- . -> reference on self
- .. -> parent reference  
for root this would just be itself..
- \$PWD -> working directory. getcwd in C
- chdir / fchdir -> cd in C

Example for getting the current working directory in C:

```
int main (int argc, char** argv) {
    char *wd = malloc (PATH_MAX);
    getcwd (wd, PATH_MAX);
    printf ("Current WD is %s", wd);
    free (wd);
    return 0;
}
```

##### 3.2.2 Paths

- **absolute path** /home/dashie/./dashie/.zshrc  
The reason for the . in the middle is that a canonical path is an absolute path, but without either .. or . in the middle.
- **relative path** ../ai-app/ai-app.tex
- **canonical path** /home/dashie/.zshrc  
can be received with "realpath".

##### 3.2.3 Max Path

POSIX systems can have different max path lengths, these are defined in "limits.h".

Macros:

- **NAME\_MAX** max length of filename (exclusive 0 termination)
- **PATH\_MAX** max length of path (inclusive 0 termination)
- **\_POSIX\_NAME\_MAX** minimal value of NAME\_MAX according to posix
- **\_POSIX\_PATH\_MAX** minimal value of PATH\_MAX according to posix

### 3.2.4 Rights

There are 3 permission categories, each having 3 groups, *read-write-execute* with *owner-group-other*

Technically there is one more bit, the sticky bit, but this is not really used anymore.

110-110-110 -> owner and group can do all, other can't do anything

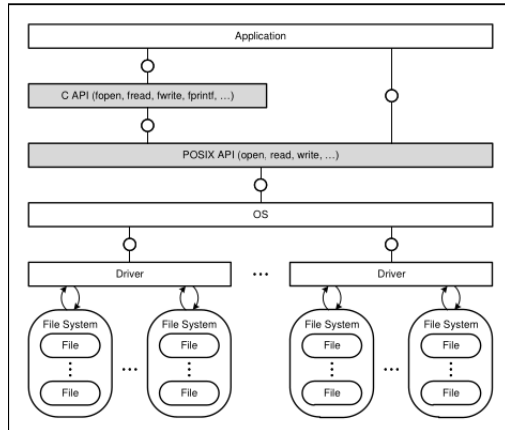
Note that this can also be done with numbers -> 7 would be all -> 111

```
drwxr-xr-x dashie dashie 884 B Fri Nov 11 11:15:56 2022 Desktop
drwxr-xr-x dashie dashie 778 B Mon Feb 20 11:03:11 2023 Documents
drwxr-xr-x dashie dashie 274 B Tue Mar 7 14:13:01 2023 Downloads
drwxr-xr-x dashie dashie 146 B Tue Dec 20 12:33:04 2022 Games
drwxr-xr-x dashie dashie 1.2 KB Mon Mar 6 13:35:10 2023 gits
```

These rights are also stored as Macros in POSIX -> "sys/stat.h"

These can be chained with | -> S\_IRWXU | S\_IRGRP

## 4 File API



The main difference with the POSIX API to the C API, is that the POSIX API gives us raw data, without any interpretation of the data, while the C API supports more specific things such as sockets, decoding etc.

### 4.1 POSIX File API

General Info:

- All file functions are declared in `<unistd.h>` and `<fcntl.h>`
- error codes can be checked with "errno"
- raw data
- should only be used for binary data, not for anything that needs interpretation

#### 4.1.1 Usage of errno

```
if (chdir("docs") < 0) { // type is int
    if (errno == EACCESS) { // EACCESS defined in the function documentation
        printf ("Error: %s\n", strerror (errno));
        // or you can use perror
        perror ("Error"); // this makes use of the standard error stream
    }
}
```

Note, not all function set this flag, and should be used immediately, as other function will overwrite it.


Codes for the error are directly defined in the function documentation.

Returns the address of a string, which describes the code in text.

perror is the same as strerror but with a special error stream.

#### 4.1.2 File-Descriptor

- valid within a process
- indexed in a table of all open files in a process
- Process file table is indexed in the global table of all open files
- Receives data in order to identify physical file (correct hardware with correct driver)
- State defined: knows current offset (offset of byte that will be read next)



In each process open file index for a process there are 3 predefined file descriptors

```
// STDIN_FILENO = 0 -> standard input
// STDOUT_FILENO = 1 -> standard output
// STDERR_FILENO = 2 -> standard error
```

#### 4.1.3 Opening files with POSIX API

```
int open (char *path, int flags, ...)
```

- O\_RDONLY: Read only
- O\_RDWR: read and write
- O\_CREAT: create file if not exists, needs another parameter for access rights
- O\_APPEND: set offset to end of file before each write access
- O\_TRUNC: set length of file to 0

#### 4.1.4 Close files with POSIX API

```
int close(int fd)
```

deallocates file descriptor fd, which can now be used by other functions.

returns 0 if ok, and -1 for error

#### 4.1.5 Usage of open and close

```
int fd = open("filename.file", O_RDONLY);
if (fd < 0) {
    // le error handling
}
// do something with file
close(df);
```

#### 4.1.6 Read data with POSIX API


```
ssize_t read (int fd, void *buffer, size_t n)
// ssize_t is a signed size_t
```

- tries to copy the next n(parameter) bytes to current offset from fd to the buffer
- returns count of read bytes or -1 when error
- blocks thread until: n bytes are copied, error occurs, end of file has been reached
- increments offset of fd by the amount of read bytes

#### 4.1.7 Write data with POSIX API

```
ssize_t write (int fd, void *buffer, size_t n)
// ssize_t is a signed size_t
```

- tries to copy the n bytes from the buffer to the offset on fd
- returns count of written bytes or -1 on error
- blocks thread until n bytes are written, error occurs, or the end of file has been reached
- increments offset of fd by the amount of bytes written



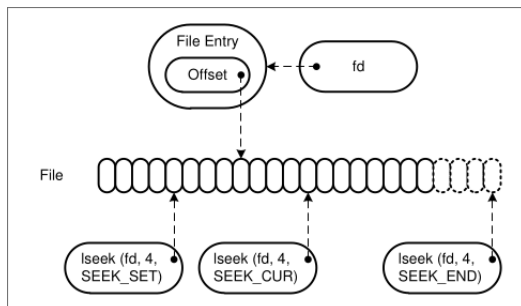
#### 4.1.8 Jump in a file with POSIX API

```
off_t lseek (int fd, off_t offset, int origin)
```

- parameter offset that should be the new offset of fd
- origin = SEEK\_SET for start of file
- origin = SEEK\_CUR for current offset
- origin = SEEK\_END for end of file
- returns new offset or -1 on error

Example usage:

```
lseek (fd, 0, SEEK_CUR) return current offset
lseek (fd, 0, SEEK_END) return size of file
lseek (fd, n, SEEK_END) go beyond end of file, which will make write put 0s in this space if called.
```



#### 4.1.9 pread and pwrite in POSIX API

```
ssize_t pread (int fd, void *buffer, size_t n, off_t offset)
ssize_t pwrite (int fd, void *buffer, size_t n, off_t offset)
```

These are alternatives to write and read, that do not change the offset, however, this means that you will need to define where we are currently!

#### 4.1.10 windoof proprietary paths

- \ instead of / because fuck you
- \ also needs to be escaped if you want to write it
- root directory per disk instead of per system
- C is the default disk, A and B were reserved for floppy disks
- functions in windoof:
  - open -> CreateFile
  - read -> ReadFile
  - write -> WriteFile
  - lseek -> SetFilePointer
  - close -> CloseHandle

## 4.1.11 Example for reading and writing data in POSIX API

```
#define N 32
char buf[N];
char spath[PATH_MAX];
char dpath[PATH_MAX];

/* get paths from somewhere */

int src = open(spath, O_RDONLY);
int dst = open(dpath, O_WRONLY | O_CREAT, S_IRWXU);
ssize_t read_bytes = read(src, buf, N);
write(dst, buf, read_bytes);
close(src);
close(dst);
```

## 4.2 C Stream API

Idea: Operating systems do things differently, even something as simple as a newline is handled differently, so we need an API that can translate this to the correct symbol:

```
// Windows: \r \n = 13d 10d = 0Dh 0Ah
// Linux: \n = 10d = 0Ah
// Mac OS: \r = 13d = 0Dh (before Mac OSX, now just like penguinOS)
```

- OS independent
- stream-based: symbol-oriented
- can be buffered or unbuffered
  - dependent on the implementation, transparent for applications
- normally buffered for files
  - independently transfers data-blocks between files and buffers
- Has a file Position indicator
  - for buffered streams: defines position in buffer
  - for unbuffered streams: is the offset in the file-descriptor

## 4.2.1 FILE datastructure

- has information about a stream
- should not be used directly, instead only per pointers that are created via the C-API
- should not be copied, pointer can be used as ID by the API
- three predefined standard-streams
  - FILE \*stdin
  - FILE \*stdout
  - FILE \*stderr

## 4.2.2 open file with C-API

```
FILE * fopen (char const *path, char const *mode)
// flags
// "r": like O_RDONLY
// "w": like O_WRONLY | O_CREAT | O_TRUNC
// "a": like O_WRONLY | O_CREAT | O_APPEND
// "r+": like O_RDWR
// "w+": like O_RDWR | O_CREAT | O_TRUNC
// "a+": like O_RDWR | O_CREAT | O_APPEND
```

- creates FILE-Object and stream for the file
- returns pointer to created FILE-Object or 0 on error

## 4.2.3 close file with C-API

```
int fclose (FILE *file)
```

- calls fflush
- closes stream defined by file parameter
- removes file from memory
- returns 0 when ok, otherwise EOF

## 4.2.4 flush of file with C-API

```
int fflush (FILE *file)
```

- writes content to write from memory into file (if content exists)
- will automatically be called when the buffer is full or file is closed
- returns 0 when ok, otherwise EOF

## 4.2.5 Conversion from POSIX-API to C-API

```
FILE * fdopen (int fd, char const *mode)
// like fopen, but instead of path, we use a file descriptor

int fileno (FILE *stream)
// returns file-descriptor for the stream, or -1 on error
```

## 4.2.6 read from file with C-API

```
int fgetc (FILE *stream)
```

- reads the next byte from stream as unsigned char and returns it as int
- increments the file-position indicator by 1

```
char * fgets (char *buf, int n, FILE *stream)
```


- reads to n-1 symbols from stream, or until newline or EOF
- adds a 0 and creates a null terminated string
- returns buffer(string) or 0 on error
- increments the file-position indicator by the amount of read symbols



## 4.2.7 "un"read from file with C-API

```
int ungetc (int c, FILE *stream)
```

- puts c(parameter -> file-descriptor) back to the stream into a special *ungetc stack*
- the next fgetc call will prefer the symbols in the ungetc stack
- no change in the file itself
- ungetc stack has a minimum size of 1 -> works at least once, can work multiple times depending on implementation
- returns c(parameter -> file-descriptor) or EOF on error



## 4.2.8 Write to file with C-API

```
int fputc(int c, FILE *stream)
```

- converts file-descriptor c in unsigned char and writes it to stream
- returns either c or EOF
- increments the file-position indicator by 1

```
int fputs (char *s, FILE *stream)
```

- writes the symbols from the string s until the 0 termination symbol into the stream
- the 0 termination will not be written
- returns EOF on error

## 4.2.9 End of file and Error in File C-API

```
int feof (FILE *stream)
// returns 0 when end of file has NOT been reached

int ferror (FILE *stream)
// returns 0 when NO error occurred

// Example usage:
int return_value = fgetc (stream);
if (return_value == EOF) {
    if (feof (stream) != 0) {
        /* EOF reached */
    } else if (ferror (stream) != 0) {
        /* error occurred, check errno */
    }
}
```

## 4.2.10 Manipulation of file-position indicator with C-API

```
long ftell (FILE *stream)
// returns the current file-indicator
// POSIX extension of ftello with return type off_t

int fseek (FILE *stream, long offset, int origin)
// set the file-position indicator like lseek
// POSIX extension of fseeko with off_t as type for the offset

int rewind (FILE *stream)
// reset the stream
// equivalent to fseek(stream, 0, SEEK_SET) and clear the error state
```

## 4.3 Ext2

## 4.4 Ext4

## 5 Processmodels

## 5.1 Base

Whenever we run a program, there are at least 2 actors: the program itself and the OS.  
 The communication between OS and program runs on the C-API (or soon rust as well :) )  
 Each program only knows itself and the OS.  
 Systems like this are called *monoprogramming*

### 5.1.1 Improvement with "async"


This is what javascript does, it is only pseudo parallel in the sense that there is *efficient process switching*, however we are still at 1 process at a time! There are multiple names for a process in this case -> task, job, or simply process.

Here we still want to ensure that each program may think that is "the only running program", this means that the operating system *must provide each service individually* -> memory, IO, etc.

## 5.2 Process

Each process has:

- Text section: image of the program in the memory (binary)
- data section: global variables of the program
- memory for heap
- memory for stack



### 5.2.1 Program vs Process

Process

- active: actually executes instructions
- executes instructions in serial manner
- may perform actions for multiple programs -> according to POSIX


Program

- passive: only says what to do
- can be run in parallel -> multiple processes execute different tasks for the program

### 5.2.2 Process Control Block (PCB)

The operating system stores information about a process in this block.

- ID for process -> PID
- Memory for the state of the processor
- scheduling information
- data for synchronization and communication between processes
- filesystem relevant information -> current open files
- Security-information



### 5.2.3 Interrupts, Processes and Context Switch

For every interrupt the Operating system saves the current state of the process in the PCB.

The following will be saved:

- Registers
- Flags
- Instruction Pointer
- MMU-Configuration -> Page-Table-Pointer

After that, the interrupt handler will be called (which is run by the OS), which can switch context if needed -> e.g. switch to another process.

As soon as the handler is done, the new context PCB is recovered.

### 5.2.4 Creation of a process

POSIX:

1. create process
2. load a program into this process and put it into ready mode


Windoof:

both at once... (proprietary)

## 5.3 Process Hierarchy

In POSIX, each process has exactly 1 parent, other than the root process.

Each process can have unlimited amount of child processes!



### 5.3.1 fork function

The fork function creates an exact copy of the parent process.

```

pid_t fork(void)
// not to mistake with :(){ :|:& };;
// kek
  
```

- the forked process has its own ID -> PID
- the forked process has the parent ID saved as well
- the return of this function will be handled in *both processes*  
in parent: return will be either the ID or -1 for error  
in child: returns 0

Example of usage:

```
int main() {
    pid_t new_pid = fork()
    // from here on both processes run this.

    if (new_pid > 0) {
        // this is the parent process
        // only the parent will run here
    } else if (new_pid == 0) {
        // this is the child process
        // only the child will run here
    } else {
        // error handling for parent process if forking failed
    }
}
```

### 5.3.2 exit function

```
void exit(int code)
```

This is a return method to get out of a process.

The code is simply the error handling code that you would like to pass to the parent process.

### 5.3.3 wait function

```
pid_t wait(int* status)
```

- blocks process until *one of his child processes ends*
- the status will cover the return code of the child process
- status usually handled with macros:  
WIFEXITED(\*status) != 0 -> if child exited properly  
WEXITSTATUS(\*status): exit code of child
- will return -1 if error  
ECHILD: no child anymore to wait for (if false you still have children)  
EINTR: was interrupted by signal


### 5.3.4 waitpid function

```
pid_t waitpid (pid_t pid, int *status, int options)
```

Like wait, but you can choose the child to wait for with pid.

- pid > 0: waits for child with this pid
- pid == -1: waits for any child -> like wait
- pid == 0 and pid < -1 enables waiting for processes of a specific process group

### 5.3.5 Examples



Worker example:

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#include <errno.h>
#include <sys/wait.h>

void spawn_worker(int a) {
    if (fork() == 0) {
        // we are in child
        printf("%d", a);
        // print number of process
        exit(0); // return code ok
    }
}

int main() {
    int pid = 0;
    for (int i = 0; i < 10; i++) {
        spawn_worker(i);
        // create 10 child processes
    }

    // do something in parent

    do {
        pid = wait(0);
        // status code is not saved!
    }

    while (pid > 0 || errno != ECHILD);
    // run until no more children!
}
```

### 5.3.6 exec functions

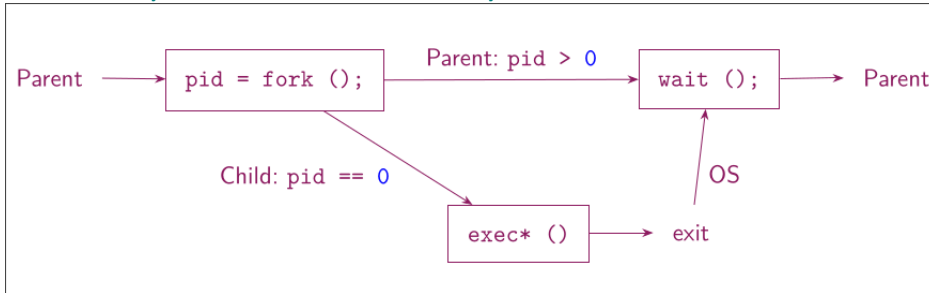
There are 6 versions of exec -> execl, execl, execlp, execv, execve, execvp

Each exec function *replaces the program in the current process with another*.

Parameters:

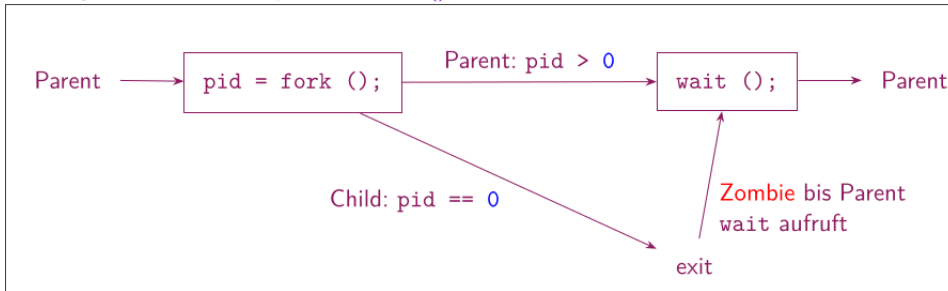
- execl\*: binary and args as list -> execl(path0, arg0, arg1, ...)
- execv\*: binary and args as vector -> execv(path, argv)

- `exec*e`: allow array for environment variables as parameter, in the other variants the env-vars stay the same
  - `exec*p`: search for filename via environment variable `PATH`, the others use absolute or relative paths
- the `*` means any character -> functions are a family



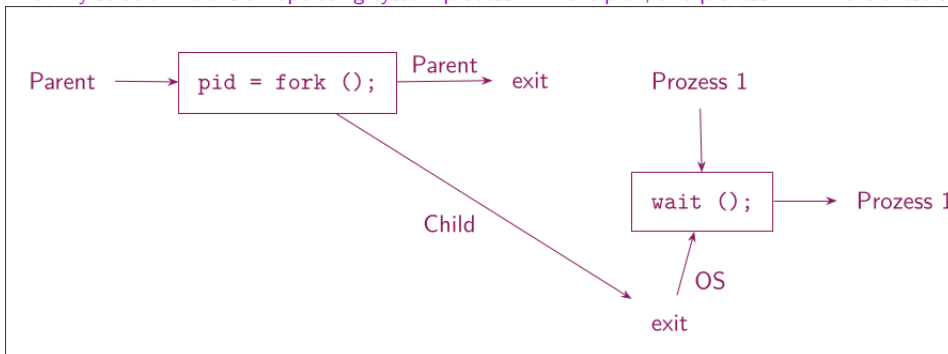
### 5.3.7 Zombie Prozesse

This refers to child processes that have ended, but aren't removed yet.  
 This means the PCB etc are all still there, but it doesn't do anything.  
 This stays like this until the parent calls `wait()` for this child!



### 5.3.8 Orphan Process

This is even worse than regular zombie processes. This means that the parent process is now dead, therefore no process can wait for this.  
 The only solution here is an operating system process with the `pid1`, this process will inherit these child processes and will then continuously end all of them.



Note, it might be that the parent process is stuck/has an error, in this case the child processes will be passed to 1 for ending them, while the parent is killed.

### 5.3.9 function sleep

```
unsigned int sleep(unsigned int seconds)
```

- waits "closely" for the amount of seconds provided
- can be interrupted by the OS (example: signals)
- returns the amount of seconds that would remain to sleep  
 example if interrupted after 5 seconds, returns entered time - 5 seconds

### 5.3.10 function atexit

Here you can pass functions that will be executed before the process exits.  
 This is useful for resources that the OS doesn't know/care about.

```
int atexit(void (*function)(void))
```

- `atexit` can be called multiple times to register multiple functions
- same function can be registered multiple times
- functions will be called in reverse order to registered order!

### 5.3.11 function to read pid

These functions simply return the pid of the current process or the parent.

```
#include <stdio.h>
#include <unistd.h>
int main() {
    pid_t my_pid = getpid();
    pid_t my_parent_pid = getppid();
    // do something with pid
}
```

## 6 C Toolchain

The toolchain handles the *precompilation, compilation, assembling and linking* of executables.

### 6.1 Preprocessor

The preprocessor translates Macros and header files into actual code.  
 The output will then be passed to the translation unit and the compiler.


### 6.2 Linker

After the individual assembly files have been assembled, the linker takes all object files and creates one executable from them.

### 6.3 Loader

The loader takes executables and dynamic libraries and loads them into ram.

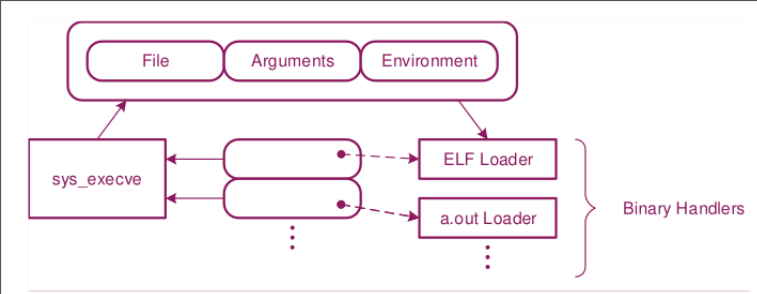
Note that static libraries are basically inside the executable.



#### 6.3.1 Penguin Loader

The loader under linux has the following steps:

1. exec\* have syscalls on sys\_execve
2. searches for file, checks for permission and opens specified file
3. counts and copies arguments and environment variables
4. provides request for each registered "binary handler"
5. Binary handler try to load data and interpret data one by one
6. program is now executed

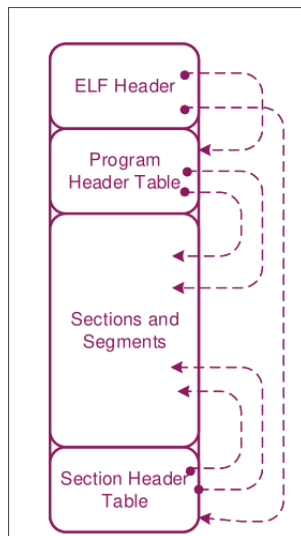


### 6.4 Executable Linking Format ELF

- Binary format, which specifies a program as binary window and penguin have some more than just this
- two possible formats
  - Linking view
  - Execution view
- Used for linker and loader
  - Object-File: Linking-view -> something.o
  - Programs: Execution-view -> no ending?
  - Shared-Objects: dynamic libraries (Linking and Execution View!) -> wroots.so
- item 4


#### 6.4.1 ELF Structure

- Header
- Program Header Table: (only necessary in execution view)
- Segments: (only necessary in execution view)
- Section: Header Table (only necessary in linking view)
- Sections: (only necessary in linking view)



#### 6.4.2 Segments and Sections

- Segments and sections occupy the same memory space
  - Compiler specifies space for sections
  - Linker specifies space for segments
- Linker combines segments with the same names and defines segments
- Compiler defines sections

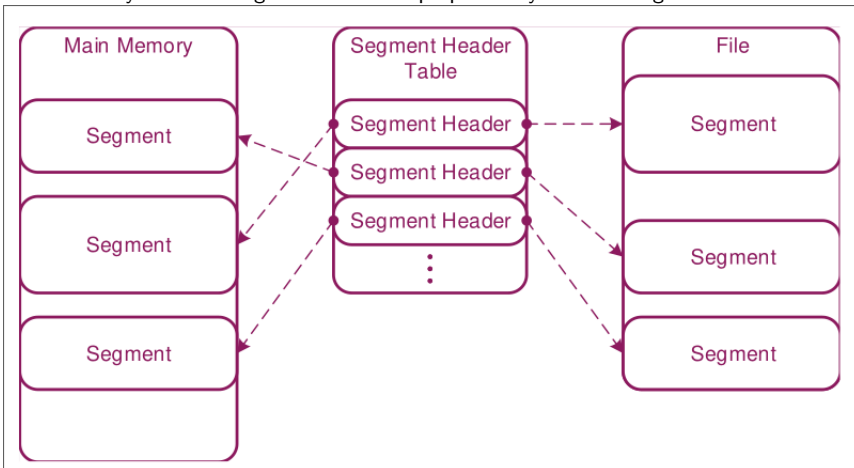


### 6.4.3 Header of an ELF file

- 52 byte: describes the structure of the file
- Type: Relocatable, executable, shared object?
- 32 or 64 bit
- Encoding: little or big endian
- Machine: i386, arm, intel 64 etc
- Entrypoint: Address at which the program should start -> main()
- relative address, count and size of the entries for the program header table
- relative address, count and size of the entries for the section header table

### 6.4.4 Program Header Table and Segments

- Program Header Table (or Segment Header Table): Table with n entries
- Each entry is 32 bit and describes a segment
  - Segment-Type and Flags
  - Offset and size of file
  - Virtual address and size in memory -> possible addition: physical address
- Segments are used at runtime by the loader
  - Loader loads specified segments to memory
  - Loader may use more segments for other purposes: dynamic linking



### 6.4.5 Sections

- Section Header Table: Table with m entries (usually :m != n | not equal sections to segments)
- Each entry (40bytes) describes a section
  - Reference to string table
  - Type and Flags
  - Offset and Size of file
  - Specific information based on section type
- Sections will be used by linker
  - collects and concatenates(same name) sections of all object files
  - generates executable

#### Section Types

- **SHT\_PROGBITS** Data defined by program, linker does not interpret this
- **SHT\_SYMTAB** Symbol Table
- **SHT\_STRTAB** String Table
- **SHT\_REL/RELA** Relocation Information
- **SHT\_HASH** Hashtable for symbols
- **SHT\_DYNAMIC** Information for dynamic Linking
- **SHT\_NOBITS** Section without data in file

#### Section Attributes


- **SHF\_WRITE** data of this section should be writable during execution
- **SHF\_ALLOC** data of this section should be in memory during execution
- **SHF\_EXECINSTR** data of this section display machine code

#### Special Sections

- **.bss**: uninitialized data
- **.data**: initialized data
- **.data1**: initialized data
- **.debug**: debug information
- **.rodata**: read only data
- **.rodata1**: read only data
- **.text**: executable instructions
- **.symtab**: Symbol-Table
- **.strtab**: String-Table

### 6.4.6 String Table

- Section of a file that has null terminated strings in a row
- String reference -> offset of string (starts at 0 and increases per char)
- typically has names of symbols
- typically does NOT include string literals like "henlo birb!", this is usually in .rodata
- item 4




#### 6.4.7 Symbols

- Symbol Table: each entry has 1 symbol
- Symbol: 16 bytes
  - Name: 4 bytes, reference in string table
  - Value: 4 bytes, depends on symbol type, could be an address for example
  - Size: 4 bytes, size of symbol (for example length of function)
  - Info: 4 bytes types, binding attributes (local, global, weak), reference to section header

#### 6.5 Static Libraries

- Static libraries are archives of object files
- will be produced with the "ar" tool
- named lib<something>.a -> note the ending
- will only be referenced if included in compilation -> clang -lmylib
- static libraries are included in the binary and will receive a static address



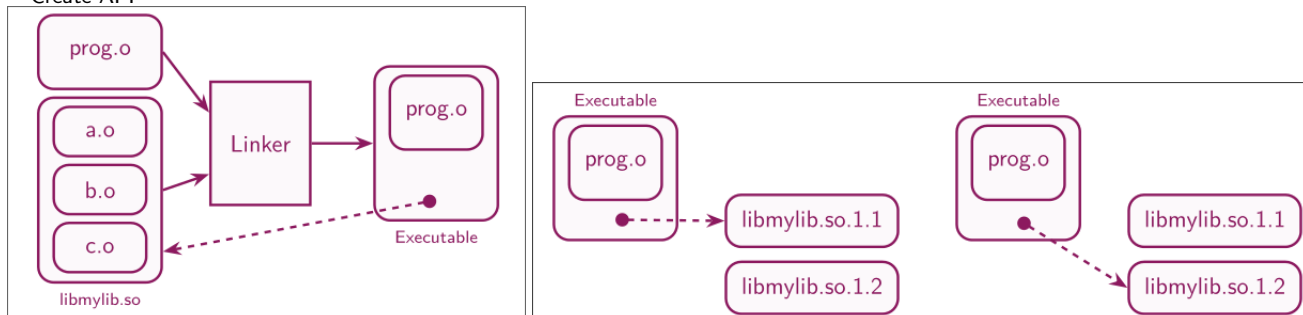
More information:

- In the early days, there were only static libraries
- Will never cause a program to suddenly stop working due to updates -> it is included in binary!
- easy to implement and use
- Must be recompiled with program if changes are made
- increases binary size
- functionality can't be increased with updates to library for binaries -> again must be recompiled
- plugins impossible or only very hard


#### 6.6 Dynamic Libraries

- loaded at runtime
- harder to implement
- executable receives only a reference to library
- library can be updated independent of binary
- if binary might stop working on library updates
- Plugins are relatively easy to create

Create API



##### 6.6.1 Delayed Loading



This is essentially lazy loading for libraries, only the libraries that will be used right then and there are loaded in order to save memory.

## 6.7 POSIX API: dynamic libraries

### 6.7.1 dlopen

opens a dynamic library and returns a handle to it

```
void* dlopen(char* filename, int mode)
```

Mode specifies how to handle the file:

- RTLD\_NOW: all symbols will be bound to the library on load
- RTLD\_LAZY: symbols will be bound if needed
- RTLD\_GLOBAL: symbols can be bound used during binding of other object-files
- RTLD\_LOCAL: symbols can't be used for other object-files

### 6.7.2 dlsym

Returns the address of a symbol from the specified library (handle for library)

```
void* dlsym(void* handle, char* name)
```

There is no type information, you only receive an address.

This means you can't know if it is a function, variable or whatever.

Example:

```
typedef int (*func_t)(int);           // our function returns int and takes int
handle = dlopen ("libmylib.so", RTLD_NOW); // open dynamic library
func_t f = dlsym (handle, "my_function"); // take function from library and store in f
int *i = dlsym (handle, "my_int");      // take variable from library and store in i
(*f)(*i);                               // call function
// Note we have no information about either f or i, we just assume they are function and variable respectively!
// You have to read API and documentation yourself to make sure that is the case!
```

### 6.7.3 dlclose

Closes the dynamic library via the specified library (handle to library)

```
int dlclose(void* handle)
```

Returns 0 on success

### 6.7.4 dLError

Returns a null terminated string as error if an error occurred.

```
char* dLError()
```

## 6.7.5 Automatic Loading of ELF files

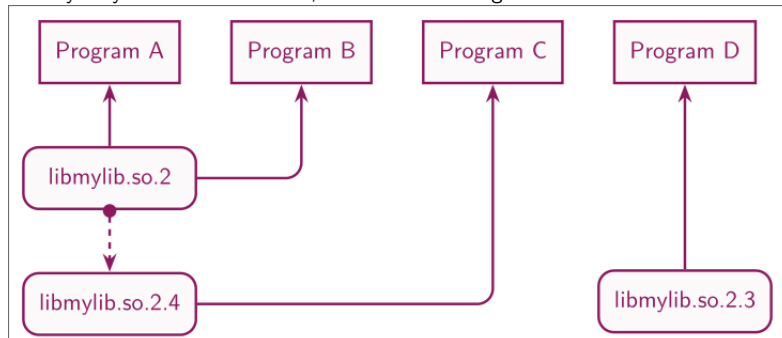
As long as there is a reference to the library in the executable ELF file, libraries will automatically be loaded by need.

## 6.8 Naming of Shared Objects

- **Linker-Name:** libmylib.so
- **SO-Name:** libmylib.so.2
- **Real-Name (filename):** libmylib.so.2.1

The tool "ldconfig" properly creates these files for you.

Usually only the realname exists, with the rest being *soft-links*




As one can see, you should always use the SO name (or linker (not recommended)) at the least, if you link against the real name, then your application will break on each update...

### 6.8.1 Updates

- Real name never changes
- SO-name: update on feature increases -> new API
- Real-Name: update on bugfixes

## 6.9 Shared Objects with Linker and Loader

Linker uses the base name, the executable the SO name and the process will finally use the real name.





## 6.10 Creating Static Libraries:

```
// compile
clang -c f1.c -o f1.o
clang -c f2.c -o f2.o

// create archive
ar r libmylib.a f1.o f2.o
```

## 6.11 Creating Dynamic Libraries:

```
// compile
clang -fPIC -c f1.c -o f1.o
clang -fPIC -c f2.c -o f2.o

// create image -> .so file
clang -shared -Wl,-soname,libmylib.so.2 -o libmylib.so.2.1 f1.o f2.o -lc
```

Important, the linker will prefer dynamic libraries, you need to override the default to use static versions if both exist

## 6.12 Using libraries


```
// static libraries
clang main.c -o main -L. -lmylib

// dynamic libraries
clang main.c -o main -lmylib

// dynamic library that will be called with dlopen
clang main.c -o main -ldl
```

### 6.12.1 ld-linux.so effective loader

- Can be used as an executable with `-list` flag
- will be called by the OS indirectly
- finds and loads all other shared objects one by one
- continues recursively to load all dependencies of loaded shared objects




## 6.13 Shared Object Facts

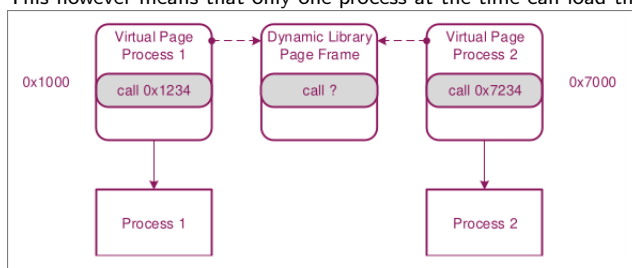
- Referenced Shared Objects are stored inside the executables
  - `readelf -d` shows the content of dynamic section
  - type of the entries is **NEEDED**
- `ldd` shows all shared objects, even indirectly used ones
  - executes executables and instructs loader to show trace
  - should only be used on trusted executables!!
- Pretty much all Executables (C) need at least two shared objects
  - `libc.so` -> standard C library
  - `ld-linux.so` -> ELF Shared Object Loader

## 6.14 Dynamic Library Implementation

Dynamic libraries must be able to be moved, this ofc since they are loaded into memory.




This however means that only one process at the time can load the library unless you load it twice, which is not something you want.



The solution to this is saving the address of the library which can be seen with 2 implementations below:

### 6.14.1 Position-Dependent Code

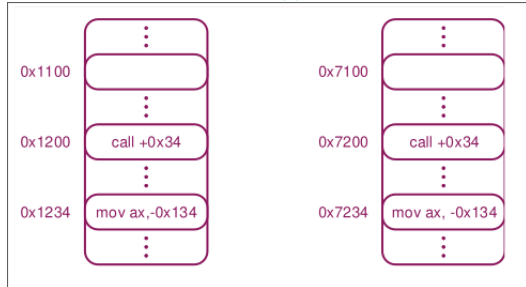
This means we jump to absolute addresses, in this case to absolute addresses of our dynamic library.



### 6.14.2 Position-Independent Code

This means we jump relatively to our current instruction pointer inside of our program.

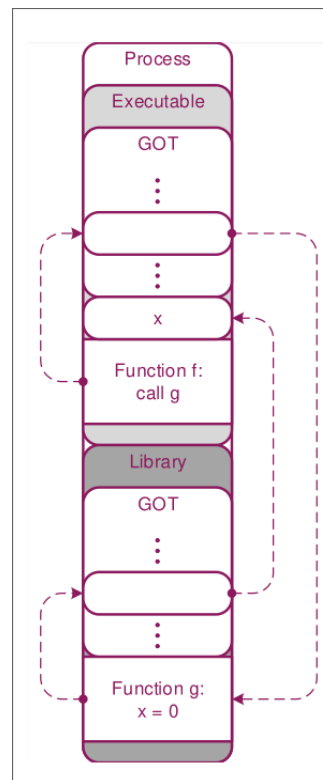
Note that the CPU has to support this, 32bit for example does not.



### 6.14.3 Global Offset Table

In each program/library we have a global offset table, which shows the relative distances to libraries etc that you need/want.

This way you can use the same dynamic library that is loaded into memory, even though another process is also using it.



### 6.15 PLT Process Linkage Table

Implements lazy binding!


- one function per entry
- PLT entry has call instruction in GOT entry
- GOT points to proxy function
- Proxy function finds link to proper function and overwrites own entry
- **Benefit:** eliminates check for "is it loaded" after initial load (is expensive!!)

## 7 MultiThreading

### 7.1 Threads


- **Threads** are parallel activities inside of a process
- **Threads** have equal access to all resources in the process
  - code(text section)
  - global variables (data section)
  - heap
  - open files
  - MMU-data

## 7.2 Process-Model vs Thread-Model




## 7.3 Thread as stack and context

- Each thread needs its own context and its own stack, since it will have its own function call stack
- Information about call stack will usually be placed in a Thread-Control-Block
- In Linux, each process has a copy of the Process-Control-Block of its own process instead!!



Example of two threads that call the same function:




## 7.4 Parallization of algorithms

Algorithms are by default serial, however, there are some parts that you could run in parallel.  
Example:

```
// this can be split
for (int i = 0; i < n; ++i) {
    sum += a[i];
}

// into this
for (int i = 0; i < n / 2; ++i) {
    sum0 += a[i];
}
for (int i = n / 2; i < n; ++i) {
    sum1 += a[i];
}
sum = sum0 + sum1;
```

### 7.4.1 Time difference



This shows the non parallizable parts on the bottom and the others on top.

As it can be seen the top parts are shrinking by large margins since they can be done in parallel!

### 7.4.2 Speedup Factor

$$f \leq \frac{T}{T'} = \frac{T}{T_s + \frac{T - T_s}{n}}$$

$$s = \frac{T_s}{T}$$

$$f \leq \frac{T}{T_s + \frac{T - T_s}{n}} = \frac{T}{s \cdot T + \frac{T - s \cdot T}{n}} = \frac{T}{s \cdot T + \frac{1 - s}{n} \cdot T}$$

$$f \leq \frac{1}{s + \frac{1 - s}{n}}$$


Legend:

- $f$ : speedup factor
- $T$ : Total time
- $T'$ : Time saved
- $T_s$ : total time that must be done in serial fashion!
- $n$ : Count of processors
- $s$ : relation of serial time to total time

Note this is often just a guess, not a calculation.

This is because we often can't say how long something would really take!

## 7.5 Amdahls Rule



$f(s,n)!!$

- best case:  $s==0$  -> everything is parallelizeable
- worst case:  $s==1$  -> nothing is parallelizeable -> user input

If you keep increasing the processor count, then you will eventually go towards the limit of 1:

$$\lim_{n \rightarrow \infty} \frac{1-s}{n} = 0$$

$$\lim_{n \rightarrow \infty} s + \frac{1-s}{n} = s$$

$$\lim_{n \rightarrow \infty} \frac{1}{s + \frac{1-s}{n}} = \frac{1}{s}$$

## 8 POSIX Threading API

## 8.1 Creating a thread

```
int pthread_create (
    pthread_t *thread_id,
    pthread_attr_t const *attributes,
    void * (*start_function) (void *),
    void *argument
)
```

- Creates a thread and returns 0 on success, otherwise errorcode
- The ID of the new thread will be stored in thread\_id
- attributes is an opaque object, with which you can specify things like the stack-size
- start\_function is the first instruction the new thread will execute
- the thread will place the parameter argument inside of that function
- the programmer must ensure the start\_function and the argument are compatible
- You must specify who is responsible for garbage collection of the data structure
- **Don't delete the stack of the thread during its execution!**

## 8.1.1 Example for pthread

```
struct T {
    int value;
};

void *my_start(void * arg) {
    struct T* p = arg;
    printf ("%d\n", p->value);
    free (arg);
    return 0;
}

void start_my_thread (void) {
    struct T* t = malloc (sizeof(struct T));
    t->value = 109;
    pthread_t tid;
    pthread_create (
        &tid,
        0, // default attributes
        &my_start,
        t
    );
}
```

The default attributes are handled by the operating system!

They are used when you pass 0 -> nullptr to this function

## 8.1.2 Attributes for pthread

Attributes are handled the following way:

```
pthread_attr_t attr; // variable for attribute
pthread_attr_init(&attr); // initialize attribute
pthread_attr_setstacksize(&attr, 1 << 16); // set the stacksize of attribute
pthread_create(..., &attr, ...); // create attribute
pthread_attr_destroy(&attr); // delete attribute
```

The number 0 means that we pass default attributes

## 8.2 pthread\_exit

```
void pthread_exit(void *return_value)
```

- ends the thread and returns the return\_value
- this is the equivalent to the return of the start\_function with a return value

### 8.2.1 Lifetime of a thread

A thread lives until:

- the thread executes a return statement
- the thread calls `pthread_exit`
- another thread calls `pthread_cancel`
- the thread is killed otherwise

### 8.3 pthread\_cancel

```
int pthread_cancel(pthread_t thread_id)
```

- Sends a demand to the OS to kill a thread with the `thread_id`
- The function doesn't wait for the thread to be killed
- The return code is 0 when the thread exists, otherwise ESRCH errorcode

### 8.4 pthread\_detach

```
int pthread_detach(pthread_t thread_id)
```

- Removes the memory which a thread has allocated, as soon as the thread is terminated
- Doesn't end the thread, in case the thread has not ended
- the return value is 0 if the thread exists, otherwise ESRCH

### 8.5 pthread\_join

```
int pthread_join(pthread_t thread_id, void **return_value)
```

- waits until the `thread_id` is dead
- Takes the return value of the thread and stores it in the `return_value` parameter
- `return_value` can be NULL -> 0, in this case it will not store the return value
- calls `pthread_detach`
- the return value is 0 if the thread exists, otherwise ESRCH

### 8.6 pthread\_self

```
pthread_t pthread_self(void)
```

Returns the ID of the current running thread

## 9 Thread-Local Storage (TLS)

The problem here is that there are certain things that are stored globally, take `errno` for example.

If this would be thread overreaching, then another thread might overwrite `errno`, therefore interfering with another thread that is running somewhere else.

This leads to race conditions where multiple threads want to access -> read/write this variable.

[TLS simply stores these variables per thread, resolving this issue.](#)

This is complicated to implement for the OS, but for many programming languages this is simply handled with an attribute.

[The idea is that you create a TLS-Variable which will be saved in a global variable as the key.](#)

[This key can then be used in the other threads with special variables. see below:](#)

### 9.1 pthread\_key\_create

```
int pthread_key_create(pthread_key_t *key, void (*destructor)(void*))
```

- creates a new key in the parameter `key`
- `pthread_key_t` is an opaque datastructure
- The OS stores each key for each thread with a value of type `void*`
- the value of this `void*` is initialized with NULL -> 0
- The OS calls the destructor for this value if it is not NULL
- returns 0 if ok, error code otherwise

### 9.2 pthread\_key\_delete

```
int pthread_key_delete(pthread_key_t key)
```

- removes a key and the associated values on all threads
- the key can't be used after calling this function
- the program must remove all memory that was allocated
- returns 0 if ok, error code otherwise

### 9.3 pthread\_setspecific and pthread\_getspecific

```
int pthread_setspecific(pthread_key_t key, const void * value)
void* pthread_getspecific(pthread_key_t key)
```

- writes or reads the value that is associated with the key for this thread
- typically used as a pointer for memory space -> []

### 9.4 Example with TLS

```
// setup
typedef struct {
    int code;
    char *message;
} error_t;

pthread_key_t error;
void set_up_error (void) {
    pthread_setspecific(error, malloc (sizeof(error_t)));
}

// read and write in the thread
void print_error (void) {
    error_t * e = pthread_getspecific(error);
    printf("Error %d: %s\n", e->code, e->message);
}

int force_error (void) {
    error_t * e = pthread_getspecific(error);
```

```
e->code = 98;
e->message = "file not found";
return -1;
}

// main and thread
void *thread_function (void *) {
    set_up_error();
    if (force_error() == -1) { print_error(); }
}

int main (int argc, char **argv) {
    pthread_key_create(&error, NULL);
    pthread_t tid;
    pthread_create(&tid, NULL, &thread_function, NULL);
    pthread_join(tid, NULL);
}
```

## 10 Communication and Synchronization

## 11 Programs and libraries

## 12 Graphical Overlays